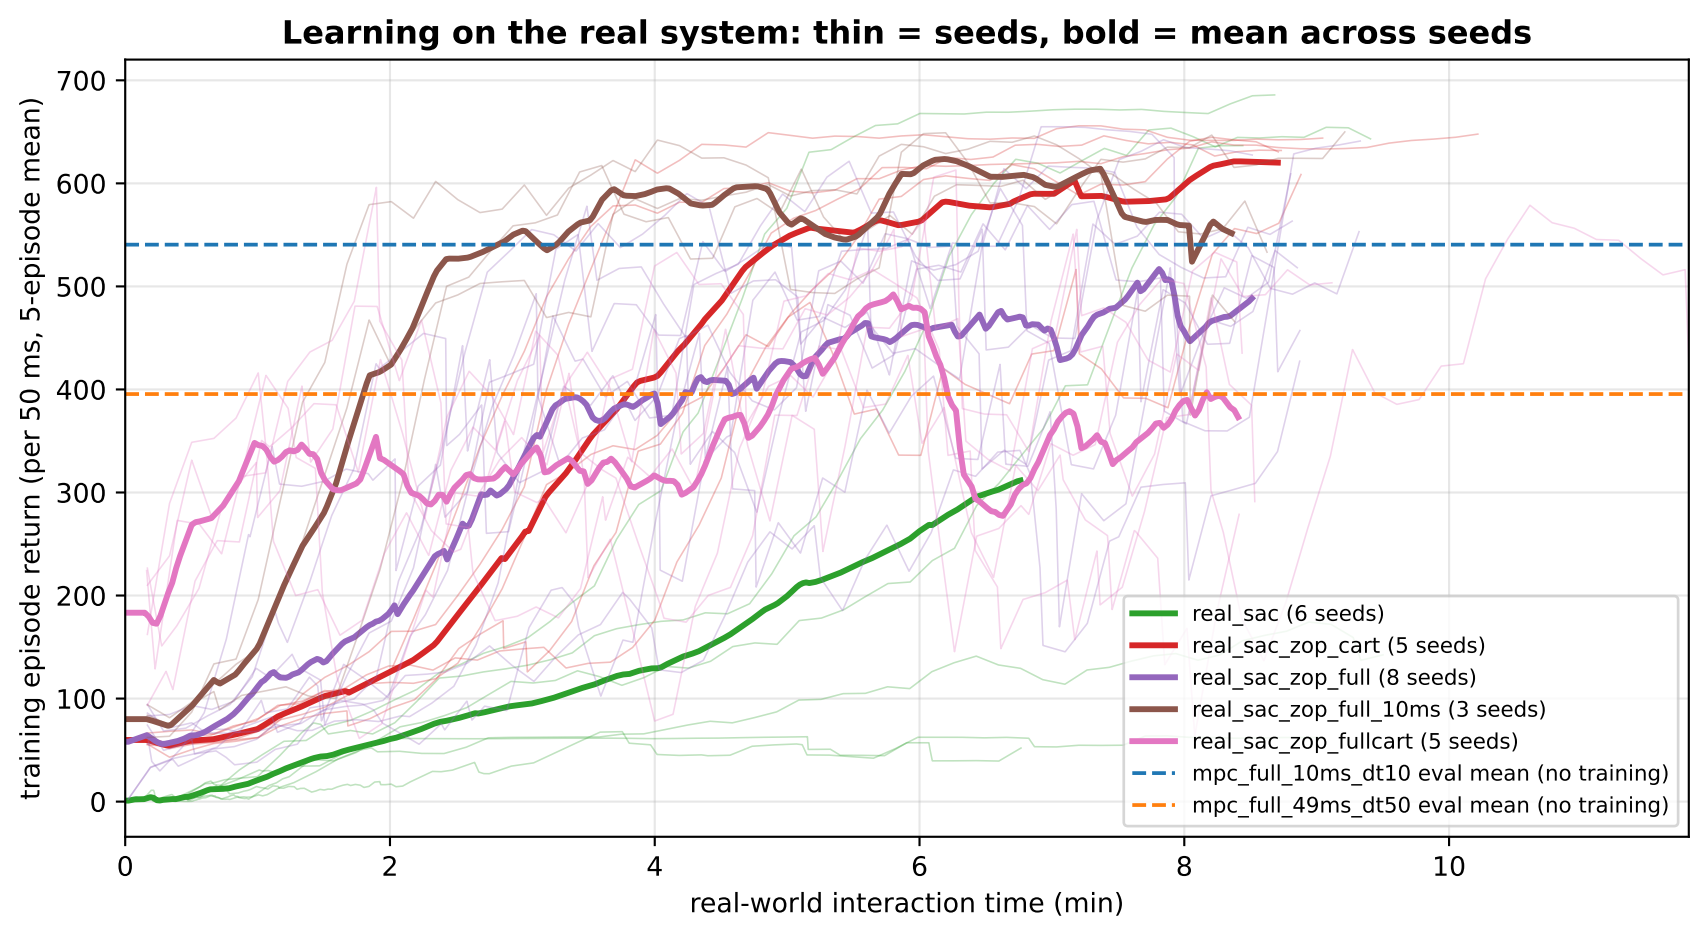
Learning on the real system: thin = seeds, bold = mean across seeds
training episode return (per 50 ms, 5-episode mean)
real-world interaction time (min)
0
100
200
300
400
500
600
700
0
2
4
6
8
10
real_sac (6 seeds)
real_sac_zop_cart (5 seeds)
real_sac_zop_full (8 seeds)
real_sac_zop_full_10ms (3 seeds)
real_sac_zop_fullcart (5 seeds)
mpc_full_10ms_dt10 eval mean (no training)
mpc_full_49ms_dt50 eval mean (no training)